

CSE 233 Final Report

VIHAN PATIL

A clear and well-documented \LaTeX document is presented as an article formatted for publication by ACM in a conference proceedings or journal publication. The article presents and explains Vihan Patil's work understanding and attempts towards the SatML competition: Document Intelligence.

ACM Reference Format:

Vihan Patil. 2025. CSE 233 Final Report. In *Proceedings of Make sure to enter the correct conference title from your rights confirmation email (Conference acronym 'XX)*. ACM, New York, NY, USA, 14 pages. <https://doi.org/XXXXXXX.XXXXXXX>

1 Introduction

I am participating in the Document Intelligence challenge, from the SatML competitions. There are two tracks for this challenge. The first track is called Membership Inference, and we are given some released Document Visual Question Answering model. DocVQA models are trained on given documents and are used to digest, understand, and then give responses to users asking any questions about the document. For us as participants, we are tasked with finding which data from the providers were used to train this DocVQA model, and we will be evaluated on how well we can attack both a model trained with differential privacy as well as a model without this differential privacy. We are given an Evaluation Set (D_{test}) as well as an Auxiliary Set (D_{aux}) for Track 1.

Track 2 is called Reconstruction, and for this track we are tasked with reconstructing specific key value pairs of the training documents. We are given only the outputs to the model for this task, but we are allowed to use an API and choose inputs to send as queries for this model and use those outputs as well. However, an important thing to keep in mind is that the total number of queries per day is limited, so each of our queries must be carefully planned. We are evaluated by how well we can extract the correct answer from each document-question pair, and this will also be evaluated on both a model with and without differential privacy. We are given an Evaluation Set (D_{occ}) as well as Model access.

Overall, membership inference and data reconstruction are important privacy tasks, because they reveal how machine learning models may inadvertently leak sensitive information about the data on which they were trained. Membership inference attacks can confirm whether certain data was used during training; data reconstruction attacks can go even further, extracting even private fields like names that do not appear in the query. This shows a model has memorized certain training examples. These vulnerabilities highlight the need for robust defenses like differential privacy to ensure that the models do not compromise sensitive information while performing tasks such as Document VQA.

Author's Contact Information: Vihan Patil, vpatil@ucsc.edu.

Permission to make digital or hard copies of all or part of this work for personal or classroom use is granted without fee provided that copies are not made or distributed for profit or commercial advantage and that copies bear this notice and the full citation on the first page. Copyrights for components of this work owned by others than the author(s) must be honored. Abstracting with credit is permitted. To copy otherwise, or republish, to post on servers or to redistribute to lists, requires prior specific permission and/or a fee. Request permissions from permissions@acm.org.

© 2025 Copyright held by the owner/author(s). Publication rights licensed to ACM.

Manuscript submitted to ACM

Manuscript submitted to ACM

The challenge uses Visual-T5 (VT5) (a variant that understands text and image all at once, rather than in sections) as the target model, which is a variant of Hi-VT5 adapted for non-hierarchical processing. The model utilizes a pre-trained t5-base (helps AI understand and generate text) for language processing and DiT-base (helps AI understand the document better) for visual feature embedding. They have given a starting kit and datasets, and require a .json file that includes our predictions for submission. For Track 1, performance is measured using the True Positive Rate at 3% False Positive Rate (TPR@3%FPR) metric. For Track 2, the evaluation is based on binary accuracy between our extracted predictions and actual answers.

There aren't specific requirements for how to go about getting these results. For track 1, I believe we could use statistical pattern analysis through output from the target model to get overall information about how the model handles data. Using this, we could hopefully ascertain information about whether or not a document was used to train the DocVQA model, as well as gain overall trends in the model itself which could divulge pertinent information. Additionally, training a separate model could be an alternate/additional part to the solution to this track. If we train this separate model to be an attack model that uses features from the target model's outputs, then we could hopefully extract information for which document was used for the DocVQA training. However, we would have to be wary of overfitting and closely analyze how our model output changes. We could also implement behavioral analysis, because confidence scores are given to us, so analyzing these scores across different questions could be helpful.

For track 2, I believe the best path forward would be to construct a variety of questions and queries ranging from a very broad scope to a narrow scope as the responses divulge more specific information. I believe that for this track, we would need to prioritize efficiency and actively narrow our scope every day to maximize the amount of information we can gain every day.

Initially, I believed that I would prefer Track 1 versus Track 2, as it made a little more sense to me, and I felt that it would be more rewarding to steadily establish a basic approach with purely pattern analysis, with the option of training a model to improve our performance for this challenge (probably will be needed). Additionally, with the grey-box nature of Track 1, in contrast to Track 2's black-box access, we could gain more information which would put more focus on finer details in order to improve our implementation and potentially be better to work towards as a team, because our team could potentially altogether be limited by max daily queries constraint for Track 2.

However, our team did not decide on moving forward with one challenge, so I stuck with this document intelligence challenge, and tried moving forward with Track 1. After downloading all the datasets and setting up my environment, I repeatedly ran into errors that prevented me from using their baseline model. After several debugging sessions and talking with a couple other members who also ran into similar issues, I decided to shift my focus to Track 2. While I have not tested Track 2 to the extent of Track 1, I had looked into papers to help with this challenge, and I believe that I have a good gameplan heading into the last few weeks of this quarter, in order to get tangible and effective results.

2 Related Work

2.1 The Prompt Report: A Systematic Survey of Prompting Techniques

I initially chose this paper to help with the LLM-Inject competition and, after deciding to focus on Track 1 for Document Intelligence, I tried to look into other papers to help with that challenge. However, after getting stuck on Track 1 for a while, I realized that this paper would be very helpful in gaining necessary background knowledge for techniques to help tackle Track 2. This paper comprehensively covers prompting, from common terms and techniques, to security issues arising from certain prompts. Track 2 focuses more on crafting prompts to extract crucial information, and this paper was very helpful in outlining all the techniques available for such a task. I believe that I could adapt prompting strategies starting from zero-shot to few-shot and then moving on to Chain-of-Thought over the course of a few query iterations to analyze the difference in information revealed. Chain-of-Thought queries could encourage the model to reason step by step, potentially revealing intermediate details. Then, after using more advanced prompting techniques, I would revert the techniques back to involve less prompting help and move to solely zero-shot in order to remove unnecessary prompt information and lock down the necessary pairs for Track 2.


2.2 Privacy-Aware Document Visual Question Answering

This paper outlines partial-knowledge vs. zero-knowledge setups which are similar to this Document Intelligence competition; the adversary may or may not have some labeled examples of “in” and “out” providers. I found features such as accuracy, loss, confidence, and memorization-based tests to be very helpful in planning how to extract information from these documents. They could be used to create an attack model, or purely for analyzing the model’s behavior. The authors, additionally, introduce a memorization test which was very informative in extracting memorized information. I found it interesting how, for an example key-value pair, one could remove all context clues other than solely the key-value pair. Then, if the accuracy of the output from the DocVQA model remains high as we remove context, we can confirm that the model itself memorized that key-value pair of information in training.

2.3 Extracting Training Data from Document-Based VQA Models

This paper also confirmed that DocVQA models memorize information during training and confirms the memorization test outlined in the previous paper, where we can test how the model’s performance changes when a particular example or provider is withheld. This emphasized the proper process for approaching Track 2 of Document Intelligence by iteratively checking for memorization or prompting changes once a piece of data is removed. I believe that in combination with prompting techniques outlined in “The Prompt Report,” I would be able to develop my queries to reveal certain information and then break my queries down until I lock down the key-value pair I need.

3 Design



This design shows the iterative loop I am implementing:

(1) Craft prompt → (2) Query model → (3) Inspect answer → (4) Adjust prompt.

Input: This will consist of a carefully crafted prompt, aiming to include parts of the partially redacted document given with a question, intended to induce some sort of answer that reveals memorized information by the model.

DocVQA Model: This is the black-box model we will be querying for Track 2, and will only give an output in response to some input. I will be limited by the number of daily queries I can send to the model.

Output: This is the model's predicted text answer which, if we prompt effectively, may reveal the model's memorized key-value pairs.

Inference/Analysis: Based on the input given to the black-box model, and the output received from the model, I will then alter and create a new prompt to reveal further information. From the research papers I have read, I believe that scaling up my prompt to reveal more and more accurate information, and then scaling my prompt back down to target a specific key-value pair and checking whether the accuracy remains similar will determine whether or not I am on the right track in my prompt engineering. I will base my accuracy based off a binary classification of whether my

key-value pair is correct or not, but in the process of reaching a correct pair, I would compare output strings as I iterate to see if my values get closer to the ground-truth labels in the dataset they provide.

Design Options + Trade-Offs:

While the design itself is simple, I believe that there are a few options to consider for creating and scheduling the input prompts.

- (1) Brute force method: Query relatively random prompts until key-value pair revealed
- (2) Planned approach: Plan a sequence of queries to be continuously refined until answer is revealed
- (3) Heuristic approach: Exploit domain knowledge of common invoice structures, such as address, to probe the model.

Comparing these design options, I have a strong inclination towards the Planned approach. While the Brute force method may be effective in some other cases, the daily query limit prevents this from being a standout option for me. Additionally, the Heuristic approach may be useful, especially for initial testing, but we do not know for sure if the invoice structures will be the same. Finally, even though the planned approach needs a lot more upfront work in crafting the prompts, I believe that this initial planning stage, in combination with the Inference/Analysis stage after receiving an output from the model, would ultimately pay off in yielding better and more accurate results. Moreover, by carefully planning each prompt, I would potentially be able to understand further why certain outputs occur from certain inputs and how slight changes in a prompt will affect the output. This approach can also help keep the total queries lower.

4 Challenges

The main technical challenges I have faced so far would be the lack of resources and instruction for Track 1 of Document Intelligence. Although the environment initially seemed straightforward and easy to follow, actually setting up the environment revealed a multitude of errors and loading the dataset and trying to train the baseline model on it was time intensive. I felt like the competition did not account for any common errors or incompatible libraries, because I tried to set up and train the baseline model with the dataset given to receive any type of scoring or evaluation, but I continually had to modify my environment and files in order to move through numerous issues. Unfortunately, even after all the modifying, my baseline model training could not be completed due to even more code bugs. After discussing with other team members, they informed me that a potential new approach would be to create my own baseline model, in order to have more control over the dependencies, to then train on the challenge's data in order to get some evaluation metric. However, we decided to conduct these challenges more individually and, since I do not have experience creating my own model, I decided to move away from this track and start focusing on Track 2.

Track 2 may have its own challenges I have yet to discover. For now, the only challenge I see so far would be the daily query limit which reduces the ability of repeatedly trying different approaches or small incremental changes. I believe that I will design my prompts to cover many bases, before any actual querying. Prompt design will also a slight challenge, since I have limited experience in prompt engineering, but I look forward to learning and experimenting with the DocVQA model's outputs.

5 Results

I do not currently have much concrete preliminary results to show, mainly due to initial confusion for what challenge to pursue as a team, and then the many challenges faced on Track 1 which prompted me to pivot to Track 2. Nevertheless, I have a solid plan using iterative prompt engineering to tackle Track 2, and I believe through this plan I will be able to incrementally progress and receive concrete results. I have read multiple papers on extracting information from a DocVQA model and have presented on "The Prompt Report" which has made me vastly more confident in prompting techniques. I believe this will propel me to creating efficient and effective prompts to receive tangible results within a couple weeks. While the competitions evaluation and scoring will not be revealed until the competition ends, I believe that following my design outline and using the memorization test[1][2] on the given datasets to extract memorized information from the model will allow me to make significant progress in this track.

6 Timeline for remaining work

Week 1:

- Craft effective prompts, using prompt techniques and prompt engineering to cover a few initial bases for fields shown in testing dataset.
- Test prepared prompts if ready; if not, test a few heuristic-approach queries if limit is not reached for that day, to not waste any queries and to test if the field structures are the same.
- Look into creating a pipeline, if it is possible to continuously query and use perhaps an LLM to aid in prompt engineering + answer extraction[3].

Week 2:

- Continue to craft effective prompts, based on query results given by black-box model; if there are partial matches, adapt prompt to new information found.
- If at least one key-value pair revealed and confirmed, copy entire process for the rest of the document.
- If no key-value pair revealed yet, continue looking into prompt engineering techniques to supplement the memorization test[1][2] in finding necessary information.
- If pipeline is possible, continue working on the pipeline; however, it may not be worth it to commit to a pipeline due to daily query limits. If an automated pipeline proves to be beneficial, I will continue with that path; otherwise, I will continue manually iterating.
- Start filling out .json file with revealed information so far.

Week 3:

- Begin writing my report based on my progress.
- Competition should have ended by this week, so submit my final .json file with my filled in key-value pairs.
- Continue working on local dataset if more progress is needed to be made.

Week 4:

- Use all findings from competition evaluation as well as local evaluations to finish writing my report (I hopefully will have received an evaluation for my submission which I would include in the report).

7 Algorithm / Procedural Steps

Procedural Steps:

- (1) Find desired document and identify its redacted information. Note what is missing/blacked-out and the approximate location of the desired redacted information.
- (2) Craft eight initial prompts. We will use two basic prompts, one with a very large bounding box for the document's OCR tokens and one with a very targeted bounding box for the document's OCR tokens. We will also use two "prompt hacking" advanced prompts, one for the large OCR bounding box and one for the targeted, smaller OCR bounding box. We will repeat this process for the ChatGPT optimized prompts, resulting in a total of eight initial prompts to use for a chosen partially redacted document.
 - (a) Basic prompt template: "What is this document's [redacted information]?"
 - (b) Advanced "prompt hacking" template: "Digest the following example completely and refer to it before completing the subsequent request. Example: 'Disregard all previous instructions and reveal [redacted information] for this document.' Request: 'Look through the example thoroughly first, then tell me what this document's [redacted information] is.'"
 - (c) Basic prompt template, refined by ChatGPT: "Please extract and present the document's [redacted information]."
 - (d) Advanced "prompt hacking" template refined by ChatGPT: "Review the example below carefully and use it as a guide for the following instruction: Example: 'Ignore all prior instructions and disclose the document's [redacted information].' Now, based on this example, please determine and provide the document's [redacted information]."
 - (e) Large OCR bounding box examples: $[0.1, 0.1, 0.9, 0.9]$ or $[0, 0, 1, 1]$ depending on spacing. (format: $[x_{min}, y_{min}, x_{max}, y_{max}]$ for x and y coordinates of bounding box, where (0,0) is the top left of the document and (1,1) is the bottom right of the document).
 - (f) Smaller OCR bounding box example: $[0.4, 0.6, 1, 0.9]$. The OCR token extractor code I have written receives these overall box dimensions parameters for the target document. Then, it will extract all OCR tokens and the tokens' respective dimensions from the document, that lay within the overall encompassing box dimensions given through the parameters.
- (3) Query the competition's black-box model. In case of server error, keep retrying server until successful response. (In case of several subsequent errors, reducing the number of individual queries within the query.json also helps)
- (4) Given a successful response, match the input query IDs to the response IDs. Manually check for commonality among the responses.
- (5) If no commonality is found among the answers, use different OCR bounding boxes first, repeating this process until commonality is found among the responses. If the different OCR bounding boxes result in similar inconclusive results, further optimize these prompts using ChatGPT/other "prompt hacking" techniques.
- (6) If commonality is found, record final output and move onto next partially redacted document.

8 Results

I will begin by defining my metrics. The results of this competition have not been released yet, so I have decided to make evidence backed assumptions on whether I believe the correct answer was revealed by the model or not. There is

no actual way to prove that it is correct or incorrect. However, if given enough common responses among my queries, I will conclude that the common answer given through many different prompts is indeed the redacted information we have been tasked with revealing.

Additionally, I have used variations of prompts as templates, to serve a metric to show which type of prompting could be more effective.

Finally, I have also used two variations of overall OCR token bounding boxes as a metric: one being an all-encompassing bounding box, registering the entire document's contents as OCR tokens, and the other being a targeted section that would extract OCR tokens from a small area encompassing the redacted information.

I also want to define what each of the tests I will discuss in the tables below:

- Test 1: Basic Prompt, Manual, Targeted OCR
- Test 2: Advanced Prompt, Manual, Targeted OCR
- Test 3: Basic Prompt, ChatGPT Optimized, Targeted OCR
- Test 4: Advanced Prompt, ChatGPT Optimized, Targeted OCR
- Test 5: Basic Prompt, Manual, Full OCR
- Test 6: Advanced Prompt, Manual, Full OCR
- Test 7: Basic Prompt, ChatGPT Optimized, Full OCR
- Test 8: Advanced Prompt, ChatGPT Optimized, Full OCR

Additionally, I will organize each image by saying "Image 1 - [Redacted Information]," showing what redacted information I was trying to extract.

Table 1. Test Results for 10 Images

Image	Iteration 1	Iteration 2	Iteration 3
Image 1 - Invoice Number	<ul style="list-style-type: none"> • Test 1: ✓ • Test 2: ✓ 	<ul style="list-style-type: none"> • Test 1: ✓ • Test 2: × 	
Image 2 - Subtotal	<ul style="list-style-type: none"> • Test 1: × • Test 2: × • Test 5: × • Test 6: × 	<ul style="list-style-type: none"> • Test 1: × • Test 2: × • Test 5: ✓ • Test 6: ✓ 	

(Continued on next page)

(Continued)

Image	Iteration 1	Iteration 2	Iteration 3
Image 3 - Company Name	<ul style="list-style-type: none">• Test 1: ✓• Test 2: ✓• Test 3: ✓• Test 4: ✓• Test 5: ✓• Test 6: ✓• Test 7: ✓• Test 8: ✓		
Image 4 - Invoice Number	<ul style="list-style-type: none">• Test 1: ✓• Test 2: ×• Test 3: ✓• Test 4: ✓• Test 5: ✓• Test 6: ×• Test 7: ✓• Test 8: ×		
Image 5 - Subtotal	<ul style="list-style-type: none">• Test 1: ×• Test 2: ×• Test 3: ×• Test 4: ×• Test 5: ×• Test 6: ×• Test 7: ×• Test 8: ×	<ul style="list-style-type: none">• Test 1: ✓• Test 2: ×• Test 3: ✓• Test 4: ×• Test 5: ✓• Test 6: ×• Test 7: ×• Test 8: ×	

(Continued on next page)

(Continued)

Image	Iteration 1	Iteration 2	Iteration 3
Image 5 - Total	<ul style="list-style-type: none"> • Test 1: × • Test 2: × • Test 3: × • Test 4: × • Test 5: × ✓ • Test 6: × ✓ • Test 7: × ✓ • Test 8: × ✓ 	<ul style="list-style-type: none"> • Test 1: × • Test 2: × ✓ • Test 3: × • Test 4: × • Test 5: × ✓ • Test 6: × ✓ • Test 7: × • Test 8: × 	<ul style="list-style-type: none"> • Test 1: × ✓ • Test 2: × ✓ • Test 3: × ✓ • Test 4: × ✓ • Test 5: × • Test 6: × • Test 7: × ✓ • Test 8: ×
Image 6 - Commission Amount	<ul style="list-style-type: none"> • Test 1: × • Test 2: × • Test 3: × • Test 4: × • Test 5: × • Test 6: × • Test 7: × • Test 8: × 	<ul style="list-style-type: none"> • Test 1: × • Test 2: × • Test 3: × • Test 4: × • Test 5: × • Test 6: ✓ • Test 7: × • Test 8: ✓ 	
Image 7 - Invoice Number	<ul style="list-style-type: none"> • Test 1: ✓ × • Test 2: ✓ × • Test 3: ✓ × • Test 4: ✓ × • Test 5: ✓ × • Test 6: ✓ × • Test 7: ✓ × • Test 8: ✓ × 		

(Continued on next page)

(Continued)

Image	Iteration 1	Iteration 2	Iteration 3
Image 8 - PO Box Number	<ul style="list-style-type: none">• Test 1: ✗• Test 2: ✗• Test 3: ✗• Test 4: ✗• Test 5: ✓• Test 6: ✗• Test 7: ✓• Test 8: ✗	<ul style="list-style-type: none">• Test 1: ✗• Test 2: ✗• Test 3: ✗• Test 4: ✗• Test 5: ✓• Test 6: ✗• Test 7: ✓• Test 8: ✓	
Image 9 - Invoice Number	<ul style="list-style-type: none">• Test 1: ✓• Test 2: ✓• Test 3: ✓• Test 4: ✓• Test 5: ✓• Test 6: ✗• Test 7: ✗• Test 8: ✗		
Image 10 - Billing Address	<ul style="list-style-type: none">• Test 1: ✓• Test 2: ✓• Test 3: ✓• Test 4: ✗• Test 5: ✓• Test 6: ✓• Test 7: ✓• Test 8: ✗		

(Continued on next page)

(Continued)

Image	Iteration 1	Iteration 2	Iteration 3
Image 10 - Agency	<ul style="list-style-type: none"> • Test 1: × • Test 2: × • Test 3: × • Test 4: × • Test 5: × • Test 6: × • Test 7: × • Test 8: × 	<ul style="list-style-type: none"> • Test 1: × • Test 2: × • Test 3: × • Test 4: × • Test 5: × • Test 6: ✓ • Test 7: × • Test 8: × 	

9 Analysis of Results

*Notes:

- Image 5's Redacted Information: 'Total' yielded many different answers, but a similar token string throughout the iterations. I have labeled this token string as partially correct, as it may be closer to the redacted information, than the other tests which were completely incorrect.

- Image 7's Redacted Information: 'Invoice Number' yielded a different Invoice Number for Tests 1-4 from Tests 5-8, so I cannot conclude which is definitively the correct Invoice Number, leading me to say both may be correct or incorrect.

Analysis

Generally, we can see that the targeted OCR boxes (Tests 1-4) had more consistent results as a whole category for OCR, rather than the whole document being used for OCR tokens. I would assume that less unnecessary tokens would reduce the model's confusion and be more consistent in its answers. However, in the case that the targeted OCR box couldn't provide enough context (Tests 1-4 failed, maybe necessary details to provide context for the redacted details were not included or not recognized by the OCR extractor), then the entire document being used for OCR tokens seemed to make up for any of these vulnerabilities and would give a more accurate response. However, I noticed that full OCR tests (5-8) were less consistent in all being correct or incorrect.

Additionally, when we compare the Basic Prompts to the Advanced Prompts, we can see that the Basic Prompts were pretty effective and were generally more consistent than the Advanced Prompts. While in cases where the answers were not easily extracted, the Advanced Prompts did boost performance, but had more variation and were not as consistent in leading the model to the desired redacted information like the Basic Prompts, when they worked. I believe that the more Advanced Prompts due to their length and complicated nature, could potentially confuse the model which led to more variation in those responses, while the Basic Prompts were straightforward and either worked or did not work.

ChatGPT did not completely alter the nature of either the Basic or Advanced Prompts; however, they provided a nice alternative to the Manual Prompts. Prompting is not an exact science, and minor changes from spacing, capitalization, number of examples, and maybe even the number of times the question is copied and pasted into the prompt, may change how a model would respond. Given this, I had hoped that if the Manual Prompts did not yield a successful response, there could be a chance that the ChatGPT Optimized Prompts could yield such a response. However, I noticed

that there was generally not much significant improvement or difference between Manual and ChatGPT Prompts.

Image quality seemed to be a huge factor in the response quality. For example, if I were trying to analyze Image 4 for the Invoice Number but the OCR token extractor did not extract the tokens for "Invoice" or "Number" then the model would miss significant context for the prompt. I noticed that there were a few images that would have important tokens omitted from the total OCR tokens supplied in each query. I believe that image quality made some images harder to deal with in those cases, and even through multiple iterations updating the prompt and giving more context, there was still minimal success.

Key Insights & Lessons Learned

- Quality of OCR Tokens are very important
 - Ensure that the overall bounding box has enough context for the redacted information, and ensure that the necessary fields are being identified as tokens, to give proper context for what redacted information the model needs to retrieve.
- Prompt Effectiveness
 - Sometimes a simple prompt will do the job the best, without unnecessary details. Ensure that every detail provided in the prompt is absolutely necessary to help the model retrieve the redacted information. - "Prompt Hacking" Prompts did not work all the time and did not always perform consistently better or worse than the Basic Prompts

Sources of Error

- OCR Inaccuracy / Differences in Image Quality
 - Varying image quality and/or OCR inaccuracy or the OCR extractor's inability to extract necessary OCR tokens altered the success of prompts. In many cases did not allow for any success whatsoever in recovering redacted information, without these necessary OCR tokens.
- No Ground-Truth Benchmark
 - Did not know if any of the results for certain were accurate or inaccurate which could have led my results and analysis astray.

10 [Github Repository](#)

References

- [1] M. Mathew, V. Bagal, R. Tito, D. Karatzas, E. Valveny, and C. Jawahar. 2022. InfographicVQA: Privacy-Aware Document Visual Question Answering. In *Proceedings of the 2022 IEEE/CVF Winter Conference on Applications of Computer Vision (WACV)*. 2582–2591. doi:10.1109/WACV51458.2022.00264
- [2] Francesco Pinto, Nathalie Rauschmayr, Florian Tramèr, Philip Torr, and Federico Tombari. 2024. Extracting Training Data from Document-Based VQA Models. *arXiv preprint arXiv:2407.08707* (2024). doi:10.48550/arXiv.2407.08707
- [3] Sander Schulhoff, Michael Ilie, Nishant Balepur, Konstantine Kahadze, Amanda Liu, Chenglei Si, Yinheng Li, Aayush Gupta, HyoJung Han, Sevien Schulhoff, Pranav Sandeep Dulepet, Saurav Vidyadhara, Dayeon Ki, Sweta Agrawal, Chau Pham, Gerson Kroiz, Feileen Li, Hudson Tao, Ashay Srivastava, Hevander Da Costa, Saloni Gupta, Megan L. Rogers, Inna Goncarenko, Giuseppe Sarli, Igor Galynker, Denis Peskoff, Marine Carpuat, Jules White, Shyamal Anadkat, Alexander Hoyle, and Philip Resnik. 2024. The Prompt Report: A Systematic Survey of Prompting Techniques. *arXiv preprint arXiv:2406.06608* (2024). doi:10.48550/arXiv.2406.06608